

Dear Hiring Manager,

I'm a software engineer who enjoys building products in places where the workflow is messy, the stakes are real, and good engineering can make the system feel dramatically simpler for the user. Most of my work has lived at that intersection of **product, systems, and automation**: taking operationally complex problems, breaking them into clean technical boundaries, and shipping software that is reliable enough for production teams to trust.

At OXmaint, I work on an AI-enabled EAM/CMMS platform built with **Next.js/TypeScript** and **Golang**, where a major part of my role has been modernizing legacy workflows into a more scalable product foundation. That has meant more than adding features: I've worked on the shape of the platform itself — redesigning admin-facing flows, improving extensibility for enterprise use cases, and shipping core operational modules that users interact with every day. Beyond engineering, I represent the product team in live demos for prospective enterprise leads and own the customer journey end-to-end — from initial technical pitch through architecture consultations tailored to each client's operational needs, to production onboarding and ongoing change-request management. This cross-functional exposure has sharpened how I think about building software: every system decision I make is informed by what I've seen customers actually need on the ground. On the implementation side, I've built **organization-scoped agentic workflows in n8n** for OxyAI, where LLM-driven actions had to be grounded in tenant context, routed through internal APIs, validated before execution, and tied into onboarding and resource-provisioning paths. I've also worked on **MQTT/HTTPS webhook-based IoT ingestion** using **Golang** and **ChirpStack**, where the challenge was not just receiving sensor data, but turning device discovery, onboarding, and downstream analytics triggers into something repeatable and operationally safe. In parallel, I've shipped enterprise integrations such as **QuickBooks** and **SAP-connected workflows**, implemented **CSP and cache-control hardening** in Next.js to strengthen frontend security, and owned parts of a **Flutter** mobile app with **MVVM, OAuth2, Firebase notifications, on-device TTS, wake-word interaction, and wearable integrations** for hands-free execution in the field.

What I like about this kind of work is that the engineering cannot be shallow. If you are building for operations teams, the product has to be intuitive, but the system beneath it also has to handle permission boundaries, retries, partial failures, integration edge cases, and real-world latency. That same pattern has shown up throughout my past roles. At Moonraft, I worked on both product-facing and data-intensive systems: scaling a **Flutter** app for a UK mortgage client to **65k+ users**, improving stability through **GA/FlutterFire profiling**, building a semantic search stack with **Spring Boot, FastAPI, Crawler4J, HuggingFace embeddings, and Elasticsearch k-NN**, and creating a React-based orchestration portal with **WebSockets, long-polling, audit logs, and alerting** to manage long-running scrape/index workflows. Earlier at Microsoft, supporting **200k+ endpoints** across **70+ enterprise customers** taught me how to think carefully about reliability, security, and customer impact in large production environments.

What I would bring to your team is a combination I've found valuable in practice: the ability to understand product workflows deeply, implement them with sound technical structure, and stay close to the operational realities that determine whether software is actually useful. I care about building systems that are not only feature-rich, but also maintainable, observable, and dependable under real usage.

I'd welcome the opportunity to speak further about how I can contribute.

Sincerely,
Sharon Dev Saseendran

Sharon Dev Saseendran

Buffalo, NY | [portfolio](#) | sharon.devs99@gmail.com | (716) 573 9185
linkedin.com/in/sharon-dev-s-8897a659/ | github.com/sharondevs

SKILLS

Languages: TypeScript · Go · Java · Dart · Python · SQL **MLOps:** PyTorch · vLLM · LangChain · Pinecone · LlamaIndex
Frameworks: Next.js · Golang · FastAPI · Flutter · N8N · Flask · Spring Boot · PostgreSQL · MSSQL · Elasticsearch · Apache Kafka · Redis
Observability: Azure (Container Apps, Functions, Service Bus, App Insights, Elastic Pool) · GCP · GitHub Actions · Prometheus · Grafana · JUnit

WORK EXPERIENCE

OXmaint Inc

Lead Software Engineer

July 2025 – Present | Sunnyvale, CA

Next.js | TypeScript | Golang | n8n | Flutter | Kafka | Redis

- Owned service architecture for OXmaint’s AI-enabled EAM/CMMS platform with a **Next.js/TypeScript** frontend and **Golang** microservices backend, SQL schema migrations, Redis caching and multi-tenant auth; migrated 12+ legacy workflows, increasing daily active usage by 40%.
- Closed 100% of audit-identified XSS vectors** by hardening the frontend with nonce-based CSP and CSRF rotation; contributed to GitHub Actions CI/CD pipeline with Docker staging, Prometheus + Grafana alerting (P95 latency, error rates), and structured logging.
- Improved multilingual onboarding** by shipping i18n-based locale routing with middleware-driven auto-locale detection and per-locale SSG.
- Reduced customer onboarding effort by 40%** by shipping 3 OxyAI agentic workflows on n8n — org-scoped context retrieval, tool-calling over internal APIs, LLM routing (GPT-OSS 120B) with automatic fallback, and Redis-backed session state with token-budget windowing.
- Cut sensor setup time by 60%** by engineering the end-to-end IoT ingestion pipeline — ChirpStack sensor networks, MQTT services (Eclipse Paho Go), automatic device discovery, tenant-scoped onboarding, and downstream analytics triggers for predictive maintenance.
- Worked on third-party enterprise integrations with QuickBooks (REST/OAuth 2.0) and SAP BTP (OData V2) — field-level mapping for work orders, purchase orders, and billable costs with automated invoice generation and PO reconciliation.
- Drove the end-to-end customer journey** — from prospect demos and technical discovery through solution architecture, onboarding, and go-live; managed CSRs to ensure **retention and expansion**.
- Owned the distribution build of the OXmaint AI Flutter app**, shipping production CMMS features with modular MVVM architecture, secure OAuth2, Firebase notifications; built on-device TTS, **wake-word detection**, and **Meta wearable (MWDAT SDK)** integration to enable hands-free task execution, **improving field-user turnaround by 30%**, expanding voice-assisted coverage across major CMMS workflows.

Moonraft Innovation Labs (acquired by UST)

February 2023 - July 2024 | Bangalore, India

Developer 1

Spring Boot | Flask | Flutter | FlutterFlow | Elasticsearch | Pinecone | ETL | MongoDB

- Delivered a cross-platform **Flutter** app for a UK mortgage client, scaling MVP to production for **65k+** users and achieving **85%** crash-free sessions via **GA** and **FlutterFire** profiling; build splash-screen transitions with **AnimatedBuilder**, **SVG masks**, **ClipPath**.
- Plug-and-play semantic-search module:** Spring Boot/FastAPI + Crawler4J(web crawling), Apache Kafka job queue monitoring, HuggingFace embeddings indexed in Elasticsearch k-NN. Delivered a React portal (WebSockets/long-polling, audit logs, alerts) orchestrating **20+ concurrent scrape/index runs**; templated configs enabled partner onboarding in **2–3 days**.
- Worked on a Product Comparison Engine catering ecommerce platforms, leading a 3-member team to orchestrate ETL pipeline to **scrape 50k+** product listings across multiple brands at scale, generating text & image embeddings to perform cosine similarity search.

Microsoft

August 2021 - February 2023 | Hyderabad, India

Cloud Solution Architect

Microsoft Intune | Azure AD | MECM | Microsoft Defender

- Managed 70+ enterprise accounts** as SME for Microsoft Intune, Azure AD, and MECM — led customer onboarding, change management, and endpoint security strategy across **200K+ managed endpoints**.
- Trained 100+ engineers** through 3–5 day workshops on Azure, Azure AD, Intune, and Defender 365; delivered technical advisory and built solution POCs for premier customers on M365 and Azure, **boosting client-side resolution rate by 15%**.

EDUCATION

University at Buffalo, State University of New York (CGPA: 3.74 / 4)

New York, United States

Masters in Computer Science (Specialization in AI/ML)

August 2024 - December 2025

Relevant Coursework: Machine Learning · Deep Learning · Data Modeling and Query Language · Computer Vision & Image Processing

Cochin University of Science and Technology, School of Engineering (CGPA: 9.36 / 10)

Kerala, India

Bachelor of Technology in Electronics and Communication

June 2017 - May 2021

PERSONAL PROJECTS

ECHO-SRE (Agentic SRE Copilot)

2026

- Built an agentic SRE copilot that root-causes production incidents via a ReAct agent over the Model Context Protocol (JSON-RPC/stdio), querying alerts, metrics, service topology, logs, and BM25-retrieved runbooks through a FastMCP server reusable by any MCP client (Claude Desktop/Cursor).
- Engineered a provider-agnostic inference gateway on the OpenAI Chat Completions API unifying local vLLM and hosted APIs (Gemini/OpenAI), with a priority-ordered chain, automatic multi-model failover (retryable-vs-fatal classification), exponential backoff, and per-provider token budgeting — removing vendor lock-in and chaining quotas (~560 req/day).
- Shipped a FastAPI Server-Sent-Events streaming service driving a React/TypeScript “SRE mode” with three per-request backends (bundled synthetic, user-uploaded scenario, or live Prometheus/Loki); instrumented with Prometheus, containerized with Docker, and CI/CD-deployed to Azure Container Apps via test-gated GitHub Actions, with a key-free offline demo and hermetic tests.

ECHO (Engaged Conversation with Human Oversight) Chat

2025

- Developed a retrieval augmented generation backend allowing Q&A based a user's profile or any documents, implementing real-time streaming responses with session-based conversation memory using Google Gemini 2.0 Flash and llamaindex.
- Built prod-ready FastAPI backend with LlamaIndex RAG framework, Google GenAI text-embedding-004 for semantic search, async document processing pipeline supporting PDF/DOCX/TXT formats and deployed with Azure Container Apps.

Visual Question Answering for Real-world accessibility

2025

- Developed a VQA system for real-world images, addressing challenges like blur, and poor lighting from the VizWiz dataset.
- We combined the pre-trained CLIP ViT-B/32 vision-language encoder and a trainable lightweight multi-head classifier to generate textual answer predictions with a test acc. of approx. **70%**.